

VWFS Machine Learning Workshop

Recommender Systems

Ahmed Rashed - UHi

University of Hildesheim (UHi), Germany

March 20, 2019

Outline

1. Introduction
2. Problem Definition
3. Dataset
4. Infrastructure and Coding Environment
5. Data Preprocessing
6. Approaches
 - 6.1 - User Average
 - 6.2 - Item Average
 - 6.3 - SVD

Outline

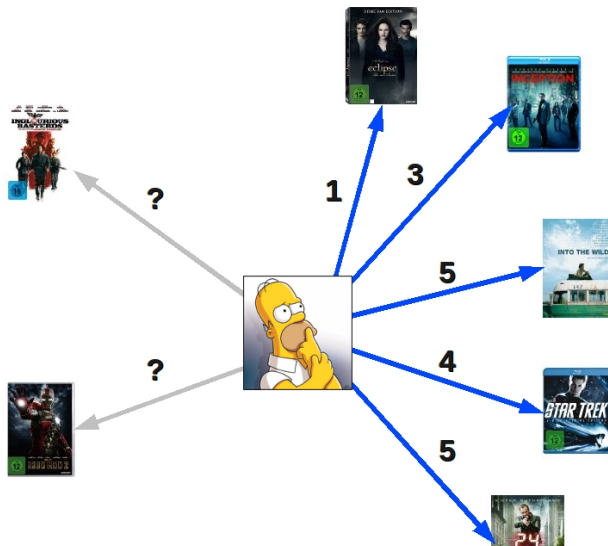
1. Introduction
2. Problem Definition
3. Dataset
4. Infrastructure and Coding Environment
5. Data Preprocessing
6. Approaches
 - 6.1 - User Average
 - 6.2 - Item Average
 - 6.3 - SVD

Why Recommender Systems?

- ▶ Powerful method for enabling users to filter large amounts of information
- ▶ Personalized recommendations can boost the revenue of an e-commerce system:
 - ▶ Amazon recommender systems
 - ▶ Netflix challenge: 1 million dollars for improving their system on 10%
- ▶ Different applications:
 - ▶ E-commerce
 - ▶ Education
 - ▶ ...

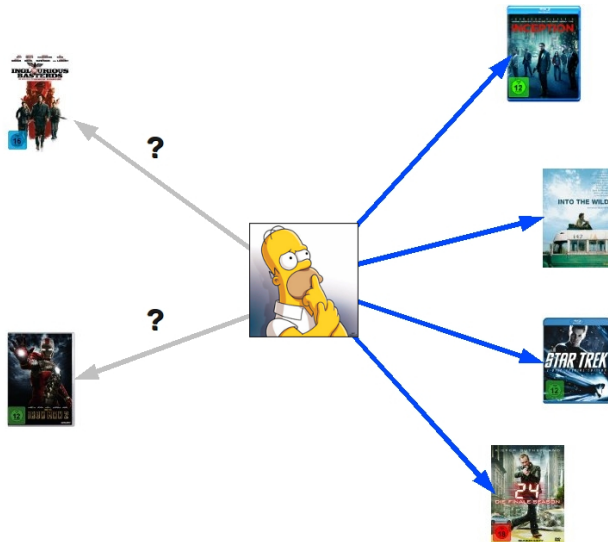
Prediction Version - Rating Prediction

Given the previously rated items, how the user will evaluate other items?



Ranking Version - Item Prediction

Which will be the next items to be consumed by a user?




Outline

1. Introduction
- 2. Problem Definition**
3. Dataset
4. Infrastructure and Coding Environment
5. Data Preprocessing
6. Approaches
 - 6.1 - User Average
 - 6.2 - Item Average
 - 6.3 - SVD

Formalization

- ▶ U - Set of Users
- ▶ I - Set of Items
- ▶ Ratings data $D \subseteq U \times I \times \mathbb{R}$

Rating data D are typically represented as a sparse matrix $\mathbf{R} \in \mathbb{R}^{|U| \times |I|}$



Recommender Systems - Task

Given a set of users U , items I and training data $D^{train} \subseteq U \times I \times \mathbb{R}$, find a function

$$\hat{r} : U \times I \rightarrow \mathbb{R}$$

That minimize the loss

$$error(\hat{r}, D^{train}) := \sum_{(u,i,r_{u,i}) \in D^{train}} \ell(r_{u,i}, \hat{r}_{u,i})$$

Mean Squared Error (MSE) is usually used as the loss function and RMSE for performance evaluation

Outline

1. Introduction
2. Problem Definition
- 3. Dataset**
4. Infrastructure and Coding Environment
5. Data Preprocessing
6. Approaches
 - 6.1 - User Average
 - 6.2 - Item Average
 - 6.3 - SVD

MovieLens 100K

- ▶ Statistics
 - ▶ 100,000 ratings
 - ▶ 943 users
 - ▶ 1682 movies
- ▶ Ratings
 - ▶ 1(Worst) to 5(Best)
- ▶ File Format
 - ▶ *userid|itemid|rating|timestamp.*

Outline

1. Introduction
2. Problem Definition
3. Dataset
- 4. Infrastructure and Coding Environment**
5. Data Preprocessing
6. Approaches
 - 6.1 - User Average
 - 6.2 - Item Average
 - 6.3 - SVD

Infrastructure and Coding Environment

1. Login to our server

- ▶ Coding Environment : `jupyter.ismll.de`
- ▶ User Name : `vw1` to `vw7`
- ▶ Password : Same as the user name

2. Create new python 3 notebook and rename it with your name

3. Import the following libraries

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import time
5 from collections import deque
6 import tensorflow as tf
7 from six import next
```

Outline

1. Introduction
2. Problem Definition
3. Dataset
4. Infrastructure and Coding Environment
5. Data Preprocessing
6. Approaches
 - 6.1 - User Average
 - 6.2 - Item Average
 - 6.3 - SVD

Data Preprocessing

- The dataset need to be loaded and splitted into train and test

```
1 def get_split_data():
2     df_train = read_process("Data/ml100k/u1.base"
3         , sep="\t")
4     df_test = read_process("Data/ml100k/u1.test"
5         , sep="\t")
6     return df_train, df_test
```

- Read Function

```
1 def read_process(filename, sep="\t"):
2     col_names = ["user", "item", "rate", "st"]
3     df = pd.read_csv(filename, sep=sep, header=None, names=col_names, engine='python')
4     df["user"] -= 1
5     df["item"] -= 1
6     for col in ("user", "item"):
7         df[col] = df[col].astype(np.int32)
8     df["rate"] = df["rate"].astype(np.float32)
9     return df
```

Outline

1. Introduction
2. Problem Definition
3. Dataset
4. Infrastructure and Coding Environment
5. Data Preprocessing
- 6. Approaches**
 - 6.1 - User Average
 - 6.2 - Item Average
 - 6.3 - SVD

6.1 User Average

- ▶ We use the average item rate as the prediction value for all test instances
- ▶ RMSE ?

6.1 User Average

- ▶ We use the average item rate as the prediction value for all test instances
- ▶ $RMSE = 1.062$

6.2 Item Average

- ▶ We use the average Item rate as the prediction value for all test instances
- ▶ RMSE ?


6.2 Item Average

- ▶ We use the average Item rate as the prediction value for all test instances
- ▶ $RMSE = 1.090$

6.3 - SVD (Factorization models)

- ▶ Each item $i \in I$ is associated with a latent feature vector $\mathbf{Q}_i \in \mathbb{R}^K$
- ▶ Each user $u \in U$ is associated with a latent feature vector $\mathbf{P}_u \in \mathbb{R}^K$
- ▶ Each entry in the original matrix can be estimated by

$$\hat{r}_{u,i} = \mathbf{P}_u^\top \mathbf{Q}_i = \sum_{k=1}^K P_{u,k} Q_{i,k}$$



6.3 - SVD (Objective Function)

Task:

$$\arg \min_{\mathbf{P}, \mathbf{Q}} \sum_{(u,i,r_{u,i}) \in D^{train}} (r_{ui} - \hat{r}_{u,i})^2 + \lambda(\|\mathbf{P}\|^2 + \|\mathbf{Q}\|^2)$$

Where:

- ▶ $\hat{r}_{u,i} := \mathbf{P}_u^\top \mathbf{Q}_i$
- ▶ D^{train} is the training data
- ▶ λ is a regularization constant

6.3 - SVD (SGD: gradients)

$$\mathcal{L} := \sum_{(u,i,r_{u,i}) \in D^{train}} (r_{ui} - \hat{r}_{u,i})^2 + \lambda(\|\mathbf{P}\|^2 + \|\mathbf{Q}\|^2) \quad (1)$$

$$\mathcal{L} := \sum_{(u,i,r_{u,i}) \in D^{train}} \mathcal{L}_{u,i} \quad (2)$$

Gradients:

$$\frac{\partial \mathcal{L}_{u,i}}{\partial \mathbf{P}_{u,k}} = -2(r_{u,i} - \hat{r}_{u,i})\mathbf{Q}_{i,k} + 2\lambda\mathbf{P}_{u,k}$$

$$\frac{\partial \mathcal{L}_{u,i}}{\partial \mathbf{Q}_{i,k}} = -2(r_{u,i} - \hat{r}_{u,i})\mathbf{P}_{u,k} + 2\lambda\mathbf{Q}_{i,k}$$

6.3 - SVD (Stochastic Gradient Descent Algorithm)

```

1: procedure LEARNLATENTFACTORS
   input:  $D^{Train}$ ,  $\lambda, \alpha$ 
2:    $(\mathbf{p}_u)_{u \in U} \sim N(0, \sigma \mathbf{I})$ 
3:    $(\mathbf{q}_i)_{i \in I} \sim N(0, \sigma \mathbf{I})$ 
4:   repeat
5:     for  $(u, i, r_{u,i}) \in D^{Train}$  do                                ▷ In a random order
6:        $\xi_{u,i} = r_{u,i} - \hat{r}_{u,i}$ 
7:       for  $k = 1, \dots, K$  do
8:          $\mathbf{P}_{u,k} \leftarrow \mathbf{P}_{u,k} + \alpha (\xi_{u,i} \mathbf{Q}_{i,k} - \lambda \mathbf{P}_{u,k})$ 
9:          $\mathbf{Q}_{i,k} \leftarrow \mathbf{Q}_{i,k} + \alpha (\xi_{u,i} \mathbf{P}_{u,k} - \lambda \mathbf{Q}_{i,k})$ 
10:      end for
11:    end for
12:  until convergence
13:  return  $\mathbf{P}, \mathbf{Q}$ 
14: end procedure

```


6.3 - SVD Implementation (Batch Generation)

► Random batches for training

```

1 class ShuffleIterator(object):
2     def __init__(self, inputs, batch_size=10):
3         self.inputs = inputs
4         self.batch_size = batch_size
5         self.num_cols = len(self.inputs)
6         self.len = len(self.inputs[0])
7         self.inputs = np.transpose(np.vstack([np.array(self.inputs[i]) for i in range(
8     def __len__(self):
9         return self.len
10    def __iter__(self):
11        return self
12    def __next__(self):
13        return self.next()
14    def next(self):
15        ids = np.random.randint(0, self.len, (self.batch_size,))
16        out = self.inputs[ids, :]
17        return [out[:, i] for i in range(self.num_cols)]

```

6.3 - SVD Implementation (Batch Generation)

- Sequentially generated batches for test data

```

1 class OneEpochIterator(ShuffleIterator):
2     def __init__(self, inputs, batch_size=10):
3         super(OneEpochIterator, self).__init__(inputs, batch_size=batch_size)
4         if batch_size > 0:
5             self.idx_group = np.array_split(np.arange(self.len), np.ceil(self.len / batch_size))
6         else:
7             self.idx_group = [np.arange(self.len)]
8         self.group_id = 0
9
10    def next(self):
11        if self.group_id >= len(self.idx_group):
12            self.group_id = 0
13            raise StopIteration
14        out = self.inputs[self.idx_group[self.group_id], :]
15        self.group_id += 1
16        return [out[:, i] for i in range(self.num_cols)]
    
```

6.3 - SVD Implementation (Model architecture)

```

1 def inference_svd(user_batch, item_batch, user_num, item_num, dim=5, device="/cpu:0")
2     with tf.device("/cpu:0"):
3         w_user = tf.get_variable("embd_user", shape=[user_num, dim],
4             initializer=tf.truncated_normal_initializer(stddev=0.02))
5         w_item = tf.get_variable("embd_item", shape=[item_num, dim],
6             initializer=tf.truncated_normal_initializer(stddev=0.02))
7         embd_user = tf.nn.embedding_lookup(w_user, user_batch)
8         embd_item = tf.nn.embedding_lookup(w_item, item_batch)
9         infer = tf.reduce_sum(tf.multiply(embd_user, embd_item), 1)
10        regularizer = tf.add(tf.nn.l2_loss(embd_user)
11            , tf.nn.l2_loss(embd_item), name="svd_regularizer")
12    return infer, regularizer
    
```

6.3 - SVD Implementation (Optimization Function)

```
1 def optimization(infer, regularizer, rate_batch, learning_rate=0.001
2 , reg=0.0, device="/cpu:0"):
3     global_step = tf.train.get_global_step()
4     assert global_step is not None
5     with tf.device(device):
6         costl2 = tf.nn.l2_loss(tf.subtract(infer, rate_batch))
7         penalty = tf.constant(reg, dtype=tf.float32, shape=[])
8         , name="l2")
9         cost = tf.add(costl2, tf.multiply(regularizer, penalty))
10        train_op = tf.train.AdamOptimizer(learning_rate).minimize(cost
11        , global_step=global_step)
12    return cost, train_op
```

6.3 - SVD Implementation (Main Function Part1)

```

1 def svd(train, test):
2     samples_per_batch = len(train) // BATCH_SIZE
3
4     iter_train = ShuffleIterator([train["user"], train["item"]
5     , train["rate"]], batch_size=BATCH_SIZE)
6     iter_test = OneEpochIterator([test["user"], test["item"]
7     , test["rate"]], batch_size=-1)
8     user_batch = tf.placeholder(tf.int32, shape=[None], name="id_user")
9     item_batch = tf.placeholder(tf.int32, shape=[None], name="id_item")
10    rate_batch = tf.placeholder(tf.float32, shape=[None])
11    infer, regularizer = inference_svd(user_batch, item_batch
12    , user_num=USER_NUM, item_num=ITEM_NUM, dim=DIM,
13    , device=DEVICE)
14    global_step = tf.contrib.framework.get_or_create_global_step()
15    _, train_op = optimization(infer, regularizer, rate_batch
16    , learning_rate=0.001, reg=0.09, device=DEVICE)
    
```

6.3 - SVD Implementation (Main Function Part2)

```

1 init_op = tf.global_variables_initializer()
2 with tf.Session() as sess:
3     sess.run(init_op)
4     print("{}_{}_{}_{}".format("epoch", "train_error", "val_error"
5         , "elapsed_time"))
6     errors = deque(maxlen=samples_per_batch)
7     start = time.time()
8     for i in range(EPOCH_MAX * samples_per_batch):
9         users, items, rates = next(iter_train)
10        _, pred_batch = sess.run([train_op, infer]
11            , feed_dict={user_batch: users,item_batch: items
12                ,rate_batch: rates})
13        pred_batch = clip(pred_batch)
14        errors.append(np.power(pred_batch - rates, 2))
15        if i % samples_per_batch == 0:
16            train_err = np.sqrt(np.mean(errors))
17            test_err2 = np.array([])
18            for users, items, rates in iter_test:
19                pred_batch = sess.run(infer,
20                    feed_dict={user_batch: users,item_batch: items})
21                pred_batch = clip(pred_batch)
22                test_err2 = np.append(test_err2
23                    , np.power(pred_batch - rates, 2))

```

6.3 - SVD Implementation (Main Function Part3)

```
1 end = time.time()
2 test_err = np.sqrt(np.mean(test_err2))
3 print("{:3d}_{:f}_{:f}_{:f}(s)".format(i // samples_per_batch
4 , train_err, test_err, end - start))
5 start = end
```

6.3 - SVD Experiment

- What is the error value if we set the dimension to 50, batch size to 1000 and max epochs to 100 ? What happens after the 60 epoch?

6.3 - SVD Experiment

- ▶ What is the error value if we set the dimension to 50, batch size to 1000 and max epochs to 100 ? What happens after the epoch 60?
 - **Over-Fitting**
 - **RMSE = 0.932**
- ▶ How to solve it?

6.3 - SVD Experiment

- ▶ What is the error value if we set the dimension to 50, batch size to 1000 and max epochs to 100 ? What happens after the 60 epoch?

Over-Fitting

- ▶ How to solve it?

With Regularization Penalty = 0.09

RMSE = 0.922

Thank You